

Joseph Elisi

Lab 108

November 6, 2019

postlab.pdf

You must use (and cite!) at least TWO additional resources for this post-lab!

https://www.tutorialspoint.com/cplusplus/cpp_passing_arrays_to_functions.htm

<https://stackoverflow.com/questions/2987876/what-does-dword-ptr-mean/2987916>

Parameter passing

When comparing the functions I made and passing in several types, the difference between passing a value vs by reference is the way they are loaded into the parameter registers.

For understanding of qword ptr I searched <https://stackoverflow.com/questions/2987876/what-does-dword-ptr-mean/2987916>

Passing two longs: By reference By value

lea rdx, [rbp-32] mov rdx, QWORD PTR [rbp-32]
lea rax, [rbp-24] mov rax, QWORD PTR [rbp-24]

Similarly passing two long pointers:

```
12 long findSumVal(long *a, long *b){
13     long sum = *a + *b;
14     return sum;
15 }
16 long findSumRef(long *a, long *b){
17     long sum = *a + *b;
18     return sum;
19 }
20 int main(){
21     long bug = 12;
22     long *a = &bug;
23     long *b = &bug;
24     long myVal = findSumVal(a, b);
25     long myRef = findSumRef(a, b);
26 }

40 mov     rax, QWORD PTR [rbp-40]
41 mov     rax, QWORD PTR [rbp-32]
42 mov     rsi, rdx
43 mov     rdi, rax
44 call    findSumVal(long*, long*)
45 mov     QWORD PTR [rbp-8], rax
46 lea     rdx, [rbp-40]
47 lea     rax, [rbp-32]
48 mov     rsi, rdx
49 mov     rdi, rax
50 call    findSumRef(long*, long*)
51 mov     QWORD PTR [rbp-16], rax
52 mov     eax, 0
53 leave
54 ret
55 _static_initialization_and_destruction_@0(int,
```

It makes sense that instead of moving the value into the parameter register, the address is loaded instead when passing by reference.

```

call Cat::Cat() [complete object construction]
lea rdx, [rbp-112]
lea rax, [rbp-64]
mov rsi, rdx
mov rdi, rax
call catLoader(Cat&, Cat&)
mov ebx, 0

```

```

lea rdx, [rbp-208]
lea rax, [rbp-112]
mov rsi, rdx
mov rdi, rax
call Cat::Cat(Cat const&)
lea rdx, [rbp-160]
lea rax, [rbp-64]
mov rsi, rdx
mov rdi, rax
call Cat::Cat(Cat const&)
lea rdx, [rbp-112]
lea rax, [rbp-64]
mov rsi, rdx
mov rdi, rax
call catLoader(Cat, Cat)
lea rax, [rbp-64]
mov rdi, rax
call Cat::~Cat() [complete object destruction]
lea rax, [rbp-112]
mov rdi, rax
call Cat::~Cat() [complete object destruction]

```

The story is a little different for objects. In both cases, value and reference, the lea command is called. On the left I have passing by reference and on the right, value. It is clear that passing by value is more expensive. Instead of just loading the effective address of the object, a bit by bit copy is made (When the default constructor “call Cat::Cat(Cat const&)” is invoked) and then the address of that copy is loaded into a parameter register.

From tutorialspoint.com, “C++ does not allow to pass an entire array as an argument to a function. However, You can pass a pointer to an array by specifying the array's name without an index.”. It is not possible to pass an array by its value, instead only the address of the first block of memory will be passed. Since I declared a long array of size 2, the program first reserves 16 spots of memory. (I struggle to understand the redundant commands, maybe clang is inefficient)

```

long arr [2] = {1, 0};
takeArr(arr);
return 0;

```

```

13 mov rbp, rsp
14 sub rsp, 16
15 mov QWORD PTR [rbp-16], 0
16 mov QWORD PTR [rbp-8], 0
17 mov QWORD PTR [rbp-16], 1
18 lea rax, [rbp-16]

```

The first element in the array is one at address [rbp-16]. This address is loaded, through command lea, into rdi, the first parameter register.

```

long arr [2] = {1, 0};
takeArr(arr);
return 0;

```

```

16 mov QWORD PTR [rbp-8]
17 mov QWORD PTR [rbp-16]
18 lea rax, [rbp-16]
19 mov rdi, rax
20 call takeArr(long*)
21 mov eax, 0

```

(I again see redundant commands, why use rax at all?) The subsequent elements in the array are accessed through arithmetic shown below.

<pre> void takeArr(long (arr){ long risky = arr[1]; long riskier = arr[2]; long riskiest = arr[5]; } int main(){ long arr [2] = {1, 0}; takeArr(arr); return 0; } </pre>	<pre> 2 push rbp 3 mov rbp, rsp 4 mov QWORD PTR [rbp-40], rdi 5 mov rax, QWORD PTR [rbp-40] 6 mov rax, QWORD PTR [rax+8] 7 mov QWORD PTR [rbp-8], rax 8 mov rax, QWORD PTR [rbp-40] 9 mov rax, QWORD PTR [rax+16] 10 mov QWORD PTR [rbp-16], rax 11 mov rax, QWORD PTR [rbp-40] 12 mov rax, QWORD PTR [rax+40] 13 mov QWORD PTR [rbp-24], rax </pre>
---	---

Arr[5] is $5 * 8 = 40$ spots in memory away from where the base of the array, which was loaded into rax, is. I added this example to show that assembly and C++ lack bounds checking. I don't know what is at that memory location but it is not the array I made.

In C++ passing by reference is absolutely different from passing by pointer. Notice how c, a cat, is treated differently in these two implementations.

```

void catLoader(Cat * c, Cat * b){
    c = new Cat;
    cout << "MEOW, i have " << c -> numLegs << endl;
}

void catLoader(Cat &c, Cat &b){
    c = new Cat;
    cout << "MEOW, i ate this many " << c.numLegs << endl;
}

```

References are implicitly dereferenced, meaning c is not the pointer to the cat, but the cat itself. Curiously, in assembly the implementations are the same for these methods.

```

catLoader(Cat*, Cat*):
    push    rbp
    mov     rbp, rsp
    sub     rsp, 16
    mov     QWORD PTR [rbp-8], rdi
    mov     QWORD PTR [rbp-16], rsi
    mov     esi, OFFSET FLAT:.LC0
    mov     edi, OFFSET FLAT:_ZSt4cout
    call    std::basic_ostream<char, std::char_traits<char> >& std::operator<< (std::char_traits<char> >(std::basic_ostream<char, std::char_traits<char> >::operator<<(long)
    mov     rdx, rax
    mov     rax, QWORD PTR [rbp-8]
    mov     rax, QWORD PTR [rax]
    mov     rsi, rax
    mov     rdi, rdx
    call    std::basic_ostream<char, std::char_traits<char> >::operator<<(long)
    mov     esi, OFFSET FLAT:_ZSt4endlcSt11char_traitsIcEERSt13basic_ostreamIT_0_E56_
    mov     rdi, rax
    call    std::basic_ostream<char, std::char_traits<char> >::operator<<(std::basic_ostream<char, std::char_traits<char> >::operator<<(long)
    nop
    leave
    ret

```

```

catLoader(Cat&, Cat&):
    push    rbp
    mov     rbp, rsp
    sub     rsp, 16
    mov     QWORD PTR [rbp-8], rdi
    mov     QWORD PTR [rbp-16], rsi
    mov     esi, OFFSET FLAT:.LC1
    mov     edi, OFFSET FLAT:_ZSt4cout
    call    std::basic_ostream<char, std::char_traits<char> >& std::operator<< <std::char_traits<char> >(std::basic_ostream
    mov     rdx, rax
    mov     rax, QWORD PTR [rbp-8]
    mov     rax, QWORD PTR [rax]
    mov     rsi, rax
    mov     rdi, rdx
    call    std::basic_ostream<char, std::char_traits<char> >::operator<<(long)
    mov     esi, OFFSET FLAT:_ZSt4endl1cSt11char_traitsIcEERSt13basic_ostreamIT_0_ES6_
    mov     rdi, rax
    call    std::basic_ostream<char, std::char_traits<char> >::operator<<(std::basic_ostream<char, std::char_traits<char> >

```

Without their headings, these functions would be impossible to tell apart. In both cases, the memory location of the cats is loaded into the parameter registers.

Objects

Spending time changing around the number of data members in my class, I noticed that the data in a class is treated just like an array. In fact it seems that objects are treated almost identically to arrays the only difference being their methods are declared in a separate part of the program, and when used they are simply called.

<pre> int sum = c -> numLegs + b -> numLegs; } void catLoader(Cat &c, Cat &b){ cout << "MEOW, i ate this many " << c.numLegs << endl; } int main(){ Cat test; Cat bruh; Cat backed; </pre>	<pre> 43 mov QWORD PTR [rbp-32], 4 44 mov QWORD PTR [rbp-24], 12 45 mov DWORD PTR [rbp-16], 9 46 mov QWORD PTR [rbp-8], 69 47 mov QWORD PTR [rbp-64], 4 48 mov QWORD PTR [rbp-56], 12 49 mov DWORD PTR [rbp-48], 9 50 mov QWORD PTR [rbp-40], 69 51 mov QWORD PTR [rbp-96], 4 52 mov QWORD PTR [rbp-88], 12 53 mov DWORD PTR [rbp-80], 9 54 mov QWORD PTR [rbp-72], 69 </pre>
---	---

Three new cats are basically three new arrays.

As you can see in catLoader, data members are accessed by their offset from the base pointer.

<pre> using namespace std; class Cat{ public: long numLegs = 4; long bruh = 12; int letsgo = 9; //long sungfs = 198; //float murica = 99.9; private: long num = 69; //long numLives = 9; //string str = "hi"; }; void catLoader(Cat * c, Cat * b){ int sum = c -> numLegs + b -> numLegs; } </pre>	<pre> 1 catLoader(Cat*, Cat*): 2 push rbp 3 mov rbp, rsp 4 mov QWORD PTR [rbp-24], rdi 5 mov QWORD PTR [rbp-32], rsi 6 mov rax, QWORD PTR [rbp-24] 7 mov rax, QWORD PTR [rax] 8 mov edx, eax 9 mov rax, QWORD PTR [rbp-32] 10 mov rax, QWORD PTR [rax] 11 add eax, edx 12 mov DWORD PTR [rbp-4], eax 13 nop 14 pop rbp 15 ret .LC0: 16 .string "MEOW, i ate this many " 17 catLoader(Cat&, Cat&): </pre>
--	--

In order for assembly to know which object it is calling, it is dependent on the address (offset from the base pointer) that was passed in. Again, to find which data member in that specific object, the offset from the beginning address of the object is used.

An example of this is that when “Cat test;” is declared, its “bruh” data member is offset by 40 from the base pointer. When it is called “test.bruh”, the assembly simply looks at what is at [rbp – 40].

It is a little more complicated when accessing data members within a subroutine. As you can see in catLoader, data members are accessed by their offset from the address of the beginning of the object. Then, the 8 byte long that we are looking for is found by rax, the register that holds the beginning address of the object.

```
1  catLoader(Cat*, Cat*):
2      push    rbp
3      mov     rbp, rsp
4      mov     QWORD PTR [rbp-24], rdi
5      mov     QWORD PTR [rbp-32], rsi
6      mov     rax, QWORD PTR [rbp-24]
7      mov     rax, QWORD PTR [rax+8]
8      mov     edx, eax
9      mov     rax, QWORD PTR [rbp-32]
10     mov     rax, QWORD PTR [rax+8]
11     add     eax, edx
12     mov     DWORD PTR [rbp-4], eax
13     nop
14     pop     rbp
15     ret
16
```

Member functions are called using the call syntax, which accepts a label of an address in the program. In my class, the label `catLoader(Cat*, Cat*)`: will store the location of that subroutine. When called, the parameter for the object in question will have it’s beginning address stored.


In the case of ‘this’, the stack pointer is the memory address stored parameter rdi that was lea from the first element in the object array.

```
mov     QWORD PTR [rbp-48], 4
mov     QWORD PTR [rbp-40], 12
mov     DWORD PTR [rbp-32], 9
mov     QWORD PTR [rbp-24], 69
lea     rax, [rbp-48]
mov     rdi, rax
```

The address is then stored right under that first element $[rbp-8]$ by an offset of 8 bytes, and is used as the 'this' pointer to reference all other data in the object that was passed into the method. In this implementation 'this' was stored in rax .

$[rax+8]$ yields 12, the data member we want.

```
catLoader(Cat*):  
    push    rbp  
    mov     rbp, rsp  
    mov     QWORD PTR [rbp-8], rdi  
    mov     rax, QWORD PTR [rbp-8]  
    mov     rax, QWORD PTR [rax+8]  
    pop     rbp  
    ret  
main:
```



Thank you for coming to my ted talk.